

Use Cases and Design of Web Applications in Decentralized Finance

Johannes Hüsters



BACHELORARBEIT

eingereicht am
Fachhochschul-Bachelorstudiengang

Software Engineering

in Hagenberg

im Januar 2021

Advisor:

DI Martin Harrer

© Copyright 2021 Johannes Hüsters

This work is published under the conditions of the Creative Commons License *Attribution-NonCommercial-NoDerivatives 4.0 International* (CC BY-NC-ND 4.0)—see <https://creativecommons.org/licenses/by-nc-nd/4.0/>.

Declaration

I hereby declare and confirm that this thesis is entirely the result of my own original work. Where other sources of information have been used, they have been indicated as such and properly acknowledged. I further declare that this or similar work has not been submitted for credit elsewhere. This printed copy is identical to the submitted electronic version.

Hagenberg, January 31, 2021

Johannes Hüsers

Abstract

Use Cases and Design of Web Applications in Decentralized Finance describes the major fields of applications of decentralized financial software. Each use case is described in detail, to create a profound understanding of the technical design and operation of each of them.

It is split up in two major parts. The first part describes the financial use cases in a more abstract and high-level manner. The second part lays down a possible implementation design, based on a very specific example of a web application that implements those uses cases. While the reference implementation is written in Solidity, a high level programming language for writing smart contracts on the Ethereum blockchain, there is a slight focus on Bitcoin and other cryptocurrencies as well. A separate chapter is dedicated to discuss criticism and potential risks of this new financial ecosystem.

Kurzfassung

Use Cases and Design of Web Applications in Decentralized Finance beschäftigt sich mit den wichtigsten Anwendungsbereichen von dezentralen Finanzanwendungen. Jeder Anwendungsbereich wird tiefgreifend beleuchtet, um ein umfangreiches technisches Verständnis für deren Funktionsweise aufzubauen.

Die Arbeit ist in zwei größere Teilbereiche gegliedert. Im ersten Teil werden die finanziellen Anwendungsbereiche auf einer abstrakten Ebene überblicksmäßig beschrieben. Der zweite Teil der Arbeit zeigt eine mögliche Implementierung der Architektur anhand eines Beispiels einer Webanwendung, die die zuvor beschriebenen Anwendungsbereiche beinhaltet. Die Referenzimplementierung ist in der statisch typisierten Programmiersprache Solidity geschrieben, die zur Implementierung von Smart Contracts auf der Ethereum Blockchain dient. Es werden aber auch Bitcoin und andere Kryptowährungen erwähnt. Ein spezielles abschließendes Kapitel zielt darauf ab, sich mit Kritik und möglichen Risiken dieses neuen finanziellen Ökosystems zu beschäftigen.

Contents

| | |
|--|-----------|
| Declaration | iv |
| Abstract | v |
| Kurzfassung | vi |
| 1 Introduction | 1 |
| 1.1 Motivation | 1 |
| 1.2 Goals | 1 |
| 1.3 Structure of the Thesis | 2 |
| 2 Fundamentals | 3 |
| 2.1 Cryptography | 3 |
| 2.1.1 Public Key Cryptography | 3 |
| 2.1.2 Elliptic Curve Cryptography | 4 |
| 2.1.3 Hashing Functions | 4 |
| 2.2 The Ethereum Blockchain | 4 |
| 2.2.1 Blockchain Fundamentals | 4 |
| 2.2.2 Clients | 5 |
| 2.2.3 Wallets | 5 |
| 2.2.4 Transactions | 6 |
| 2.3 Smart Contracts | 6 |
| 2.4 Decentralized Finance | 7 |
| 2.4.1 State of the Art | 7 |
| 3 Use Cases of Decentralized Finance | 8 |
| 3.1 Classification of Financial Services | 8 |
| 3.2 Store of Value | 9 |
| 3.2.1 Proof of Work | 9 |
| 3.2.2 Proof of Stake | 9 |
| 3.2.3 Switching the Consensus Algorithm | 10 |
| 3.3 Payments | 10 |
| 3.3.1 Bitcoin | 10 |
| 3.3.2 Centralized Stable Coins | 11 |
| 3.3.3 Dynamic Supply Protocols | 11 |
| 3.3.4 Stability Protocols | 11 |

| | | |
|----------|--|-----------|
| 3.4 | Lending & Borrowing | 12 |
| 3.5 | Exchanging | 12 |
| 3.5.1 | Characteristics of Decentralized Exchanges | 12 |
| 3.5.2 | Liquidity Pools | 13 |
| 3.5.3 | Risks | 14 |
| 3.6 | Investing | 15 |
| 3.6.1 | Tokens | 15 |
| 3.6.2 | Synthetic Assets | 15 |
| 4 | Design and Architecture | 16 |
| 4.1 | Technology Stack | 16 |
| 4.1.1 | Solidity | 16 |
| 4.1.2 | Truffle | 16 |
| 4.1.3 | Ganache | 16 |
| 4.1.4 | Web3 | 17 |
| 4.1.5 | Mocha & Chai | 17 |
| 4.1.6 | React | 17 |
| 4.1.7 | TypeScript | 17 |
| 4.2 | The SafeMath Library | 17 |
| 4.3 | The ERC20 Token Standard | 18 |
| 4.3.1 | Example Implementation | 18 |
| 4.4 | Liquidity Pool Implementation | 20 |
| 4.5 | Exchange Implementation | 22 |
| 4.6 | Testing Smart Contracts | 23 |
| 4.7 | Connecting Smart Contracts to Web Applications | 24 |
| 4.8 | Smart Contract Deployment | 25 |
| 5 | Closing Remarks | 26 |
| 5.1 | Criticism | 26 |
| 5.2 | Risks | 27 |
| 5.3 | Prospective Impact | 28 |
| | References | 29 |
| | Literature | 29 |
| | Software | 30 |
| | Online sources | 31 |

Chapter 1

Introduction

1.1 Motivation

It is Thursday, January 28th, 2021. The stock market is currently experiencing very tempestuous times. The online community *WallStreetBets* (WSB) on the social media platform *Reddit* managed to create a tremendously strong hype across several securities such as *GameStop*, rising the price up to 1000% and more in just a few days. All these securities were known as common stocks that are being sold short by big hedge funds. Due to the rapid price increase, a lot of big funds had to liquidate their short positions and lost several \$10 billion dollars ¹.

The American broker *Robinhood* is known to be a common choice as a trading platform for members of the WSB community. In order to prevent the big hedge funds from even more losses, *Robinhood* decided to manipulate the market by preventing its users from buying the affected stocks. However, selling those stocks was still possible. This resulted in an enormous price drop and the company had to deal with a lot of criticism on several social media platforms.

This is just one of many examples of how a centralized financial authority abused their power in favor of the enrichment of themselves or of their allies. It happens on a day to day basis and is a big flaw of the traditional financial system. Events like the *GameStop short squeeze* led people to the creation of a new financial ecosystem, which is not characterized by censorship, geo blocking, single point of failure and intransparency. This is possible by making use of cryptocurrencies in order to create a world wide decentralized financial network to allow people to be more independent and regain control over their own finances. While the fundamentals of such a financial economy are completely different, the ways how to use financial products did not change. Decentralized applications enable the possibility to implement those use cases on a blockchain-based network, opening entire new opportunities for the financial industry.

1.2 Goals

The goal of this thesis is to provide a better understanding of the term *decentralized finance*, its most popular use cases and how those can be implemented on the Ethereum

¹known as the *GameStop short squeeze*

platform using the programming language *Solidity* and a web frontend. As Ethereum is considered as a general purpose blockchain, there is an infinite number of ideas which could be built on Ethereum. However, this thesis only focuses on the biggest financial use cases on the example of a decentralized web application (dapp).

1.3 Structure of the Thesis

Chapter 2 provides a brief introduction to the fundamentals of blockchain technology and Ethereum in particular. The major use cases of the financial industry are described in chapter 3 and how they operate specifically on decentralized platforms. The chapter 4 explains how an application based on Ethereum could be implemented that features the use cases from the previous chapter before giving a short overview on criticism, risks and a forecast in chapter 5.

Chapter 2

Fundamentals

2.1 Cryptography

The term *cryptocurrency* already suggests, that blockchains rely heavily on cryptography. But cryptography is an even broader topic and exceeds by far the ranges of blockchains, cryptocurrencies and decentralization. Without cryptography, our digital life would be very different today. Cryptography is the theory of data encryption, which means “secret writing” in Greek. Together with cryptanalysis, which is the opposite (decoding data), cryptographic can be conflated as cryptology. Cryptography makes it possible to make data content unreadable for specific people, detect unauthorized data manipulation and guarantee the authenticity of a communication partner. What cryptography cannot achieve, is prevention of unauthorized data manipulation and the prevention that data is being read at all [10, pp. 10-11].

2.1.1 Public Key Cryptography

Public key cryptography are specific asymmetrical encryption methods that utilize both private and public keys. While symmetrical encryption methods such as the *Data Encryption Standard* (DES) and *Advanced Encryption Standard* (AES) use a single key to encrypt and decrypt a message, asymmetrical encryption methods use the public key of the recipient to decrypt the message and a private key to encrypt it. At the first glance, it may seem that asymmetrical encryption methods are way better than symmetrical encryption methods and that they should be used all the time. But that is not the case. Asymmetrical encryption methods are less secure, more complex, more difficult to invent and slower than symmetrical encryption methods [10, p. 172]. Each type of encryption method has its own use cases so they actually complement each other. While symmetrical encryption methods are really good at encrypting data content, asymmetrical encryption methods are used to exchange the symmetrical keys and prove ownership and authenticity through digital signatures.

The most popular asymmetric encryption methods are *RSA*¹, *Diffie-Hellman* and *elliptic curve cryptography*. RSA and Diffie-Hellman are the most common algorithms outside of the blockchain ecosystem and are used for example in banking, telecommu-

¹acronym for its founders Rivest, Shamir, Adelman

nications and e-commerce. Because those are not used by blockchain technologies, we will focus on elliptic curve cryptography.

2.1.2 Elliptic Curve Cryptography

Elliptic curve cryptography is an asymmetric cryptographic algorithm, which is a lot more difficult to compute than RSA and Diffie-Hellman. Both Bitcoin and Ethereum use the *secp256k1* algorithm, which uses a 256 bit length [6]. An elliptic curve is the visualization of a prime order, which is nothing else than a set of points in a two dimensional coordinate system. This makes it possible to calculate a third point, by connecting two random points on the elliptic curve with each other. This works every time, except when the two points are completely vertical to each other. By multiplying a randomly generated private key with a specific point on the elliptic curve, the generator point, it is possible to create all points on the elliptic curve in an unforeseen way. The result of such a calculation is the public key, consisting of a prefix, the x and the y coordinate.

2.1.3 Hashing Functions

A hashing function maps a string of arbitrary size to a value of fixed size. But not all hashing functions are suitable for cryptographic problems. A good hashing function aims to have as few collisions as possible. A collision occurs when two different input strings lead to the same output hash. In order to prevent that as good as possible, each hash value should occur equally and a minimal modification of the input string should lead to a completely new hash value. This is important because the more collisions occur, the more vulnerable is the hashing function for hackers. If a hacker finds a string that generates the same hash value he can apply the real signature to his faked text [10, pp. 200-201].

2.2 The Ethereum Blockchain

2.2.1 Blockchain Fundamentals

Ethereum is an unbounded state machine, which consists of a globally, uniform state and a virtual machine that makes changes in that state. Blockchains like Ethereum consist of several components. The foundation is a decentralized peer-to-peer network around the world, where everyone can participate. The Ethereum main network is the second largest peer-to-peer network². Altering the state is done on a blockchain through transactions, which play a key role in every blockchain. Furthermore, a set of rules needs to be applied to the network in order to reach consensus.

There are currently two major consensus algorithms, *proof-of-work* (PoW) and *proof-of-stake* (PoS), which will be discussed in section 3.2. The global state machine, which is in Ethereum's case the *ethereum virtual machine* (EVM), uses that consensus algorithm, to process the transactions according to the defined rules. After processing them, they are grouped together to specific blocks and chained after each other in a public

²after the main network of Bitcoin

ledger. The compliance of those rules would not be possible without a game theoretical incentivization scheme. Of course it is impossible to enforce those rules as strictly as in any kind of centralized system. But the technology can be designed in such a way, that it becomes economically unfeasible to harm the network in a critical way. For example, no one can prevent you from winning guaranteed the lottery. But the system is designed in such a way that it would cost a lot more buying all lottery tickets in contrast to the anticipated win.

So far, these are all theoretical concepts. In order to become a real blockchain, there needs to be one or more open-source implementation of that protocol. The network can be as decentralized as possible, if not everyone has the opportunity to access the source code, there is still a little bit of trust involved, which should be minimized in any type of decentralized system.

2.2.2 Clients

Clients are exactly those open-source implementations, which are the primary interface to the peer-to-peer network. On Ethereum, there are a lot of different implementations written in many different programming languages. The most popular clients are *Parity* [23], written in *Rust*, and *Geth* [21] which is written in *Go*. Each client is able to perform certain tasks on the blockchain, depending on the needs of the user. Almost every client serves as a wallet consisting of addresses and private and public keys in order to have access to specific digital assets.

Clients, which keep track of the entire blockchain are called full nodes. Full nodes have certain hardware conditions because the entire blockchain can consume a lot of memory space. They play an important role to secure the network and make it more decentralized. Some other implementations can participate on the consensus algorithm to evaluate the next block, using either mining (PoW) or staking (PoS). While each of these node types can be combined in an arbitrary order, almost every node consists at least of a basic routing functionality to connect to other nodes.

2.2.3 Wallets

Wallets are the main interfaces for an average user to interact with a blockchain. They come in a lot of different forms, varying of usability, security, supported platforms, architecture and more. A common approach is to categorize them by their platforms. There can be web applications such as *MyEtherWallet*, browser extensions such as *MetaMask*, desktop, mobile or cross-platform applications. These types of wallets are all considered as *hot wallets*, since they are constantly connected to the internet. While this is very convenient for the user, it comes with a certain security risk, as the private key could be stolen by other people. *Cold wallets*, however, aim to keep the private key disconnected from the internet as good as possible. This could be done using small devices called hard wallets, where the private key never leaves the device, or through paper wallets, where the user writes down his private key on a piece of paper and stores that in a secure place.

The management of the private keys can be done in two ways. The older and easier approach is by storing the keys in a nondeterministic way, which means generating each private key independently from each other based on a cryptographic random number

generator. Nondeterministic wallets have certain disadvantages and should not be used anymore today, except for simple tests.

The more modern approach are deterministic wallets, where all private keys are derived from a seed through the use of a one-way hash function. They are easy to backup, import and export. The seed is usually derived from mnemonic code words, which makes it possible to build up a hierarchical tree structure with child keys. The different branches could be logically organized for company departments or account categories. The public keys can then be generated on an untrusted server without having access to the private keys.

2.2.4 Transactions

Transactions are the heart of the blockchain as they are the actual data which is stored in a data block. It is important to understand, that there are not real units of a coin stored on the blockchain. There are just discrete outputs, stored as integer values, which give people who own the private key the right to spend them. Transactions consist of certain inputs and certain outputs, transferring the ownership of those atomic units of value. When someone says: “A wallet received Ether”, he just means that there is a new transaction output that can be spent with the private key of that wallet. Therefore, the balance of a wallet is the sum of all transaction outputs for that wallet. After a transaction is signed with a private key, it will be submitted into the mempool, where they wait to be included into a block. The first transaction of a block is a special transaction, known as the *coinbase transaction*. It does not consume a transaction output but pays new minted coins to the miners or stakers.

But those are not the only incentives for miners or stakers. Each transaction can also include some transaction fees which are very important for the blockchain for several reasons. As mentioned before, fees compensate miners or stakers and serve as a security mechanism, making it economically infeasible for attackers to flood the network with transactions. As the fees are collected by that miner or staker who verifies the block which records that transaction, he will prefer transactions with higher fees. Therefore, transactions with higher fees are always processed faster than transactions with none or low fees. Depending on the current market situation, it might be possible that the fees are so low, that the transaction will never be included into a block. In order to prevent that, fee estimation services calculate the appropriate fee, based on capacity and the fees offered by competing transactions. Fees are not stored directly in a transaction, as every additional storage costs money. Fees are the implicit difference between all transaction inputs and outputs.

2.3 Smart Contracts

Smart contracts are programs that are executed by the *ethereum virtual machine* (EVM). There are several programming languages that are specifically designed to program smart contracts such as *Serpent*, *Vyper* and *Bamboo*, but *Solidity* is by far the most common. Smart contracts have some characteristics that are different from other programs. First of all, they only run when they are called by a transaction. A transaction can be only initiated by *externally owned accounts* (EOAs) but both contract addresses

and EOAs are able to accept transactions. Transactions can have either a value, data or both of them. When the data payload is set, it is very likely that this transaction is addressed to call a function from a smart contract.

In contrast to Bitcoin, Ethereum is turing-complete, so the programming languages give the developer all kinds of freedom to do what they want, like they are used to in other high level programming languages. While this sounds great in the first place, it is very vulnerable to introduce errors in the code. When dealing with applications in the area of *decentralized finance*, this is the last thing you want. It is not even possible to fix the errors because smart contracts are immutable. Once deployed to a blockchain, it stays there forever. The execution on the EVM is deterministic, so every user, no matter on which machine, gets the same output. Because smart contracts are turing-complete it would be possible to write code that never terminates, which costs resources of the network. In order to prevent that, the sender has to send gas fees, which pay for the execution of the smart contract. Once, the contract is out of gas, the EVM stops executing the code. Because it is very hard to predict how long turing-complete programs will run, the sender sets a gas limit on the transaction, stating the maximum amount of gas he is willing to pay. Any unspent gas is being transferred back to the sender.

2.4 Decentralized Finance

The financial system is cumbersome. Traditional finance in particular has still problems to adapt to the digital age. Financial institutions are slow, bureaucratic and based on a lot of trust. New online companies try to solve that problem with more automated software instead of trusted people, but they are still confined to the boundaries of the ecosystem. *PayPal*, for example, tried to create its own digital currency but the government of the United States of America prohibited that.

Decentralized finance (DeFi) brings the typical use cases of traditional finance into the digital world in a decentralized way. The goal is that nobody has to trust each other but still can do business with anyone. It should be completely without censorship, so it is open for everyone, no matter which country in the world they live in. And it should be more robust in terms of availability and security as there is no central party that would risk an outage of the entire system.

2.4.1 State of the Art

The term and concept of decentralized finance is still very young, many people refer to it as an experimental form of finance. It started to arouse interest in 2019 by a few technic enthusiastic people and became really popular in the world of digital assets in the middle of 2020. As of January 2021, there are approximately \$26.2 billion locked in smart contracts on the Ethereum network [32], which is by far the most popular blockchain for decentralized finance. Despite the tremendously high market dominance of Ethereum, DeFi is not confined to a single blockchain. There are already new blockchains arising that are specifically designed for use cases in the decentralized finance area like DeFiChain [31].

Chapter 3

Use Cases of Decentralized Finance

This chapter aims to introduce the five most relevant use cases of decentralized finance. Each use case is built on top of the previous one and rises in complexity. For example, storing value in digital systems is quite easy nowadays but moving real physical assets such as gold or real estate into a blockchain is still pretty demanding. Note, that each type of financial service is not specific to DeFi, which means that they are applicable to every financial environment.

3.1 Classification of Financial Services

The digitalization changed the financial industry. A few years ago, a lot of new companies with a significant online presence¹, entered the market and changed the way how people interact with financial services. And now, the industry is experiencing again a new shift to more and more decentralized platforms and applications. But even though the interaction interface changes from time to time, the underlying types of services we use in the financial sector stay the same. The fundamental service that is necessary to build a financial ecosystem is some kind of value storage. Once that is created, it is logical to have the ability to transfer this value to other people in order to compensate them for a product or service. The most popular and dominating use case of banks, which is lending, builds on top of that. Once the ecosystem gets more mature, new types of value storage are introduced and people need a way to exchange two value storages for each other.

The last and the most sophisticated use case occurs, when people try to map complex real world examples in that financial ecosystem, such as securities, insurances and derivatives. In 2021, decentralized finance is about to enter exactly that use case. On a blockchain, this is called tokenization. A token is a new value store that is built on top of another cryptocurrency. At the time of writing this in January 2021, Chainlink [18] is the token with the highest market capitalization of \$9 Billion [39] with the purpose of connecting real world data with the Ethereum blockchain².

¹often referred as fintech

²known as oracle networks

3.2 Store of Value

No matter of the financial product or service, they all need one essential thing. In fact, our entire economic system depends on it: a currency, or in simpler terms, something to store value. Storing value is a fundamental trait of money alongside with being exchangeable and having a unit of account. The last two characteristics are comparatively easy to accomplish in the online world. Storing value in a digitalized way without having a central trusted authority, however, was a similar challenge as the people had to face in the early days when they started to switch from bartering to a real currency. A common consensus needs to be created where everyone can easily verify that a specific piece of money is authentic and was not created illegally. Even further, each person that uses this money needs a proof that it is storing real value. This had not been possible until Satoshi Nakamoto, an anonymous person or group, published a specification [9] in 2008, on how digital money could be implemented. The first cryptocurrency was founded: Bitcoin.

3.2.1 Proof of Work

Bitcoin uses a mechanism called proof-of-work in order to give each block its value. This utilizes a cost-function which is designed to be easily verifiable but quite expensive to compute [3]. In order to create a new block, a value needs to be found that matches with the correct number of zero bits of the block hash. Once the correct solution has been found, the new block represents real value, because operating this computationally intensive task on a CPU costs a lot of electricity.

The integrity of this block is guaranteed as well, because in order to change the history of the blockchain, each changed block needs to be re-computed. The longest chain of blocks on the network is accepted as the *truth*, so if dishonest people want to cheat and change the history of the blockchain, they would need to create the longest chain of blocks which can be only achieved by having more than 50% of the computing power of the whole network. While these so-called 51% attacks are a threat on blockchains with a smaller number of network members, it is very unlikely to happen on a well established network like the Bitcoin blockchain [12]. The concept of proof-of-work in order to achieve consensus in a decentralized manner became very popular and can be used in other areas of applications too, such as elections, lotteries, asset registries, digital notarization and more [1].

3.2.2 Proof of Stake

A different approach to storing digital value is proof-of-stake. Instead of making the participants of the network solve computational expensive problems, the consensus is created by proving that someone owns specific funds. A blockchain that uses a proof-of-stake algorithm to achieve consensus has a set of validators. These validators are running a master node which gives them the opportunity to vote. In order to do so, the validators need to prove that they own a specific amount of funds. This is done by sending a special transaction, which locks away their currencies for a specific time. If the validators are honest and their vote corresponds to the result of the majority, they will get their funds back and an additional reward proportional to their deposited stake.

If their vote gets rejected, the dishonest validator risks losing his money [2].

3.2.3 Switching the Consensus Algorithm

At the time of writing this in January 2021, the Ethereum blockchain uses a similar proof-of-work algorithm like Bitcoin, called *Ethash*. However, Ethereum is about to switch to the proof-of-stake algorithm *Casper* in the near future [33]. Algorithms that are based on proof-of-stake have the big advantage of less energy consumption because they are not focused on computational intensive tasks. But proof-of-stake comes with a caveat as well. Implementing incentives on a voting system which are based on the amount each validator is willing to stake means that the rich validators get richer and the poor validators stay poor.

The process of changing a consensus algorithm of a blockchain is not easy to accomplish. Because the network is decentralized, there is no single entity which can force all members to change the algorithm from proof-of-work to proof-of-stake. Just letting the people choose what they prefer to use is also not a good idea, because it is very unlikely that all agree on one type of algorithm. That would probably trigger a hard fork of the blockchain resulting in two separate networks, similar to what happened to Bitcoin in August 2017.

Bitcoin's network is by design very slow in verifying transactions. A block which is mined approximately every 10 minutes has a size of about 1 MB. Due to the small block size, Bitcoin is only capable of processing 7 transactions per second. Some people wished for a Bitcoin network which is more suitable for day to day payments. That is why *Bitcoin Cash* was created, emerged out of a hard fork of the Bitcoin network [8, 29]. In order to prevent a hard fork on Ethereum due to the transition from proof-of-work to proof-of-stake, the Ethereum protocol has a built-in mechanism called *difficulty bomb*³ which makes it more difficult over time to be profitable with mining³ by increasing the size of the problem to solve [14]. This ensures a smooth transition of all members to switch to proof-of-stake.

For the majority of users on the blockchain which are neither miners nor validators it does not matter which consensus algorithm is being used. In fact, it is not even important when developing decentralized applications based on smart contracts like it is done in chapter 4. The only thing that makes an impact is, that the blockchain is able to store real value on the network by reaching consensus by its members.

3.3 Payments

3.3.1 Bitcoin

When thinking about a cryptocurrency that can be used for payments, usually Bitcoin is the one that comes first into mind. And that is justified. Bitcoin was the first cryptocurrency and has by far the highest market capitalization. But when it comes to payments, Bitcoin is a really poor choice.

Bitcoin is too volatile. It is not uncommon that the price of Bitcoin fluctuates a few \$1000 per day. Merchants who try to sell their products with Bitcoin have a really hard

³used synonymously to proof-of-work

time if the revenue of yesterday is only worth the half today. There are a lot of reasons for Bitcoin's volatility and that will not change in the near future.

Bitcoin is too slow. The block time on the Bitcoin blockchain is approximately 10 minutes. Transactions take a long time to be verified by the network. Too long for a cashier to wait for a proof that his customer made a correct transaction.

Bitcoin is too deflationary. By design, Bitcoin has a fixed supply of 21 million. As Bitcoin gets more popular over time, it is only logical that its price increases. While this is a pleasant thing for investors, it is an unwanted side effect for a financial economy, because everyone refuses to spend their coins as they increase in value.

3.3.2 Centralized Stable Coins

There are a few strategies to remove the volatility from a digital asset. The most popular way is by pegging a cryptocurrency or token to a reference fiat currency like the US dollar. This is achieved by collateralizing every unit of the digital asset with a unit in the reference currency. This works pretty well, but the problem is, that this is a quite centralized approach, as the collateralization is managed by a central authority. However, this is a very common way and the biggest stable coins such as *Tether* [38], *USD Coin* [41] and *Gemini Dollar* [34] operate using collateralization.

3.3.3 Dynamic Supply Protocols

The price of an asset is driven by supply and demand. If the supply is larger than the demand, the price of this asset falls and vice versa. Most assets have a fixed supply, or at least a fixed distribution schedule, which gives the market forces full control over the price. But if the supply is being controlled, the price can be controlled as well. Dynamic supply protocols do this by increasing the supply if the demand increases and take out supply if the market is experiencing a reduced demand. This results in an almost stable price. A token that was one of the first that implemented such a dynamic supply protocol is *Ampleforth* [16]. As the demand increases, the amount of coins in each wallet will increase as well, while keeping the same proportions of the market share.

3.3.4 Stability Protocols

Although dynamic supply protocols work pretty well, it can be really confusing for users, if the amount of that asset in their wallet fluctuates on a daily basis. Stability protocols take a different approach by controlling the price through incentives instead of changing the supply. If the demand increases, the protocol makes it less attractive to buy that asset and rewards users that buy when the demand is low. The bigger the market capitalization of that asset, the more accurate is the price. Decentralized exchanges work in a similar way as described later in section 3.5.2. A common token using stability protocols is *Maker* [36], which is by far the project with the most value locked⁴ in decentralized finance.

⁴\$4.7 Billion (as of January 2021, <https://defipulse.com/>)

3.4 Lending & Borrowing

Lending and borrowing is the main business model of banks in traditional finance. In order to assess their risks, banks force their customers to do an identification. Therefore, it is a lot easier to evaluate the probability to receive their funds back. And even if they don't, they can initiate legal ways to get their money back. In decentralized finance, this is a lot more difficult. All users are anonymous and not regulated in any way. The network needs to find new mechanisms to make sure that the users behave as expected. Section 3.2 introduces two approaches to make it more attractive for users to play by the rules: incentives and punishments. The system of incentivization is used in the consensus protocol *proof-of-work*, where people receive a reward if they help to support the network. A punishment is used in *proof-of-stake*, as people will lose their collateral, if they vote for the wrong blocks. This system is used as well in lending applications. In order to borrow a specific asset, the user needs to deposit a collateral first, before he can withdraw the loan. If specific conditions are met, e.g. the price of the collateralized asset falls below a certain value, the user will get liquidated and loses his collateral.

While this system is the most popular right now, it only makes sense for users if they think that their collateral will grow in value as time goes by, otherwise they would just swap their collateral in the asset they need. Most cryptocurrencies are deflationary by nature, so they are supposed to increase in value. Stablecoins, however, such as *Tether* or *Dai*, are designed to have a constant value as good as possible. Therefore, collateralized lending works best with a deflationary asset and a stablecoin. This is exactly what *Maker* [36], the most popular lending project, is doing using Ether and *Dai*.

As decentralized exchanges become more popular, an alternative type of lending arises: liquidity mining. It is completely different than collateralized lending and serves a real use case by providing liquidity to an exchange. Section 3.5.2 describes this technique in more detail.

3.5 Exchanging

3.5.1 Characteristics of Decentralized Exchanges

With the arise of many new fintech enterprises in the last few years, online exchanges started to grow in importance on the market as well. While they all were very centralized in the first place, DeFi made it possible to establish new ways of exchanging money. Although each online exchange may behave very different, they can be all categorized into three big types when it comes to their structure: Centralized, decentralized and non-custodial exchanges. Centralized exchanges (CEXes) are the traditional approach where all the power is centralized to one specific organization. Decentralized exchanges (DEXes), however, are the complete opposite where there is no single entity in control and decisions are made completely based on smart contracts. Non-custodial exchanges are somewhere inbetween but are often confused with decentralized exchanges. It is important to note, that none of the types are superior in comparison to the others. Each type of exchange has its own benefits and risks. In order to assess whether an exchange is suitable for a specific use case, it is crucial to know how to classify it.

Probably the most apparent indicator is the ownership of the private keys. On

CEXes, the assets are coffered for the user. If someone buys Bitcoin on a centralized exchange, private keys are never an issue. On non-custodial and decentralized exchanges you have to manage your private keys by yourself. While this is usually a good thing, because you do not have to trust someone else⁵, it comes also with the risk of losing all your assets if you forget your private key, since nobody can restore it for you.

Owning your private keys is a good indicator but by far not the only trait a decentralized exchange needs to have. When looking at the technical infrastructure of the exchange, many things could be structured in a centralized manner. While the code of a centralized exchange is usually proprietary, a DEX needs to have code that is licensed exclusively with an open-source licence, making it possible for anyone to fork the project in case the service is no longer available. That is also highly related to emergencies. How does an exchange react to bad or unexpected events such as hacking attacks? CEXes have the option to put their service temporary offline until the issue is resolved, in order to save the funds of the users. On decentralized exchanges, there is no such thing and assets may be lost forever.

Another crucial topic to consider is censorship and geo blocking. If the service is not available in the whole world, it is not a DEX. Centralized exchanges always decide which coins are listed on the exchange, which is usually a pay to play model. Decentralized exchanges do not have such regulation. Anyone can add a new pair of coins to swap on a DEX, which often leads to a tremendously high amount of unknown and irrelevant listed coins.

The last key indicator on how to categorize an online exchange is the way how the order matching and the settlement works. Centralized exchanges use the same algorithm as stock exchanges, which utilizes an order book. Buy and sell orders are listed in a centralized ledger and the settlement happens when two orders match. Non-custodial exchanges usually use price oracles which try to get price information [15, p. 47]. Truly decentralized exchanges even go a step further and solve this problem without a third party price oracle. A DEX uses a concept that relies on *liquidity pools*, which will be discussed in detail in the next section.

3.5.2 Liquidity Pools

Every decentralized exchange needs liquidity. Because there is no order book, assets have to be already on the exchange before the user wants to swap two coins. In order to achieve that, the DEX depends on liquidity providers, which add their funds to the exchange. Liquidity providers get rewarded⁶ by the DEX because nobody provides liquidity for free. The rewards can be taken from the exchange fees, which are usually still lower than on CEXes. For example, *Uniswap*, one of the biggest decentralized exchanges, rewards each liquidity provider with 0.3% of each transaction taking place on that liquidity pool proportional to their share of the pool [40]. If there is a yearly transaction volume on the pair Ether/Tether USD of \$17.7 billion and the liquidity provider owns a 0.1% share of this pool, he will be rewarded with \$53.100 per year. The share is usually expressed through liquidity tokens, which can be traded as well.

Liquidity providers can only provide liquidity in pairs. If someone wants to provide

⁵often referred to as “not your keys, not your coins”

⁶known as *liquidity mining*

liquidity to the WBTC-ETH pair, which is *Wrapped Bitcoin*⁷ and Ether, he needs to add the same value both in WBTC and ETH. People who want to swap their WBTC to ETH add only one asset of that pair and receive the other one, which creates an imbalance in that specific pool. The total value stays the same but the proportion of the assets change. In this example, there is more Ether compared to *Wrapped Bitcoin* after the swap. This makes it very attractive for the market to swap ETH back to WBTC, because the user will get proportional more WBTC for less ETH⁸. In high-volume markets this imbalance will be exploited very soon and the original proportion is restored.

If a liquidity pool of a specific trading pair is very small, there is a risk of clearing the entire pool by a single transaction, which would give away the associated asset almost for free in the next transaction. In order to prevent that, a user never gets the entire value in the swapped asset. The difference between the expected and the actual value is called slippage. That might seem annoying in the first place, but it is a crucial instrument for the exchange to work properly. This happens on exchanges with an order book too, because it is very unlikely that both a buy and a sell order have the exact same price. The deeper⁹ the order book or the liquidity pool, the lower the percentage of the price slippage. If someone makes a high volume order on a small order book or liquidity pool, he will experience high slippage because his order wipes out a lot of opposite orders on a CEX and uses a high pool percentage on a DEX [35].

3.5.3 Risks

While providing liquidity to an exchange pool might look very profitable at first, there are also certain risks the liquidity provider has to deal with. Smart contract risk is something which you are automatically exposed to when trusting software instead of people. Because software is still written by people and people make mistakes, there is no guarantee that the smart contract on a DEX works completely the way it should be. On a turing-complete network such as Ethereum with programming languages like Solidity the risk is even higher because you can literally do anything in a smart contract. Because a smart contract is deployed only once, it is impossible to fix errors once it is public on a blockchain. This risk is far less present on networks which are not turing-complete such as Bitcoin.

Even if the smart contract risk is eliminated as good as possible by writing a lot of tests, there is still the project risk, that the developers might deliberately implement errors and backdoors into their code in order to disburse liquidity to themselves. This is often the case when the project has poor auditing and if the founders are not known to the public.

Impermanent loss is the last and most present risk liquidity providers are exposed to. The explanation of a liquidity pool in the last section is missing an important factor. Even if there is no one trading an asset pair, an imbalance can still occur because each asset is volatile by itself unless the asset pair consists of two stable coins pegged to the same fiat currency. Impermanent loss always exists if one of the two assets has large

⁷Bitcoin as a token on the Ethereum network

⁸this phenomenon is called arbitrage and applies to all efficient markets

⁹in terms of more orders and more liquidity

price fluctuations in a short amount of time. And while it is usually good to diversify an investment, impermanent loss has an even greater impact on non-correlating assets, because the price difference can be much higher. The only way to reduce the risk of impermanent loss is to have a larger time horizon where you can choose a good time to cash out the liquidity. Unfortunately there is no such thing as impermanent win because it is always a loss once one of the two assets fluctuates. The loss is expressed through the loss in the liquidity pool compared to the price gains by holding each asset separately.

3.6 Investing

3.6.1 Tokens

Investing in cryptocurrencies is already exciting. But Ethereum offers an even more customizable investing experience through *tokens*. Based on different standards, applications can create tokenized products according to their needs. The most popular token standard on the Ethereum network is the ERC-20 standard, which is described in detail in section 4.3. This standard has its flaws, for example, the gas fees are very high because there are separate transactions for token approval and transfer. There are already better alternatives out there, such as the standards ERC-223, ERC-777 and ERC-820 which are built on top of ERC-20. However, they are rarely supported by wallets and applications.

The ERC-721 standard is different than the others. It makes it possible to create non-fungible tokens, where different tokens can hold different values. This feature is the foundation to tokenize real-world assets, such as stocks, real estate and works of art.

3.6.2 Synthetic Assets

While it is possible, to tokenize almost everything, from basketball players to entire supply chains of companies, the most obvious idea is to tokenize other financial asset classes such as marketable securities. There are a few good reasons, why it is likely that synthetic assets are the first major use case of tokenization. Firstly, it is not very difficult to implement. Tokens which represent another asset, just need to be backed up with the same amount denoted in that asset, similar to how stablecoins like Tether work. Secondly, synthetic assets solve the centralization issue in other asset classes, as stocks, bonds or REITs are managed by a central investing firm. Investors who hold shares of tokenized stocks, are independent from brokers, stock exchanges, business hours, regulations and more.

Chapter 4

Design and Architecture

This chapter builds on top of the concepts and use cases of the previous chapters and lays out the architecture and design of a typical decentralized application by using a more practical approach. After a short introduction into the used technologies there will be a detailed view on each important part of a modern decentralized web application. To keep this chapter more concise, only selected code snippets will be shown to describe certain aspects. However, the entire code base of the working application can be found on *GitHub* [19].

4.1 Technology Stack

4.1.1 Solidity

The main components of decentralized web applications are smart contracts. Smart contracts can be written in a few different programming languages such as *LLL*, *Serpent*, *Solidity*, *Vyper* and *Bamboo*. Solidity is by far the most popular and currently the de-facto standard in writing smart contracts [2]. Solidity is statically typed, has a JavaScript-like syntax and supports various popular programming concepts such as libraries and inheritance [11]. Because there are no classes in Solidity, it is considered not as an object-oriented but as a contract-oriented language [25].

4.1.2 Truffle

Truffle is a development environment which helps to develop decentralized applications on the EVM [26]. It handles all the difficult parts starting from smart contract compilation, testing, linking, deployment and continuous integration. When developing more than just some simple Smart Contracts, Truffle is an important tool to save a lot of time.

4.1.3 Ganache

Because testing a decentralized application on the main Ethereum network would not be a good idea (and is expensive too), some sort of testing blockchain is needed. *Ganache* is a local blockchain for developers which provides everything to test a decentralized

applications on the local machine [20]. It offers accounts with fake ether and works really well with Truffle, which is built by the same developers.

4.1.4 Web3

In order to link smart contracts to a graphical user interface such as a web client application, web3.js comes into play. It is a JavaScript library and provides all the necessary APIs in order to connect a blockchain backend to a web frontend [28]. It also manages connections to specific wallet providers such as *MetaMask*.

4.1.5 Mocha & Chai

Testing plays in blockchain development an even bigger role than in software development in general, because deployed smart contracts are immutable and bugs literally cost money. A flawed function which causes an endless loop would burn the entire gas that was available for that function call without doing anything useful. *Mocha* is a JavaScript testing framework which makes it possible to test the functionality of each smart contract in a structured and automated way [22]. Together with the assertion library *Chai* [17] and the integration into Truffle, Mocha releases its full potential as a fully fledged blockchain testing environment.

4.1.6 React

What would be the best smart contract without a graphical user interfaces which gives people the opportunity to interact with the blockchain in an easy and intuitive way. *React* is a JavaScript library for building user interfaces for the web and by far the most popular web framework out there [24]. React is super fast because it uses a virtual DOM and provides the user a modern and intuitive experience.

4.1.7 TypeScript

While other frontend frameworks such as Angular come with TypeScript out of the box, you have to add additional dependencies in React yourself. When it comes to larger projects or whenever bugs could become expensive, it is a good idea to introduce some kind of type checking in JavaScript. TypeScript is a superset of JavaScript and helps to detect errors before they occur at runtime [27]. In financial applications like this, there is usually an even higher focus on reliable, bug-free code.

4.2 The SafeMath Library

Arithmetic operations in Solidity wrap on overflow. This can lead to bugs, because most programmers assume, that it would raise an exception when an overflow occurs, as it is done in other high level programming languages. By introducing the library **SafeMath**, an entire class of bugs is being eliminated so it should be used everywhere possible.

Program 4.1: The SafeMath library.

```
1 library SafeMath {
2   function sub(uint256 _a, uint256 _b) internal pure returns (uint256) {
3       assert(_b <= _a);
4       return _a - _b;
5   }
6
7   function add(uint256 _a, uint256 _b) internal pure returns (uint256) {
8       uint256 c = _a + _b;
9       assert(c >= _a);
10      return c;
11  }
12 }
```

4.3 The ERC20 Token Standard

The ERC20 standard makes it possible to build standardized tokens on the Ethereum protocol which can be integrated into the existing ecosystem [13]. This is achieved through a common interface each token needs to implement. The standard was initially introduced in 2015 as an *Ethereum Request for Comments* (ERC). It received its name due to the automatic issue assignment on Github, which had the issue number 20 [2, 4].

4.3.1 Example Implementation

The function `totalSupply` returns the maximum amount available of this token in *wei*, the smallest unit possible. The supply can be either static or dynamic. In this case, it is static assigned through the constructor when the smart contract is created. The return type is a 256 bit unsigned integer which is the most common data type in Solidity.

```
1 uint256 private _totalSupply;
2
3 function totalSupply() public view returns (uint256) {
4     return _totalSupply;
5 }
```

Each token is responsible to keep track of all addresses holding that token. This can be stored using the data type `mapping` which is a key value pair of addresses and token amounts. The function `balanceOf` returns the token balance for a specific address.

```
1 mapping(address => uint256) private balances;
2
3 function balanceOf(address _owner) public view returns (uint256 balance) {
4     return balances[_owner];
5 }
```

The functionality of sending tokens to a specific address is implemented in the `transfer` function. It takes an address and an amount as parameters and returns a boolean if the transaction was successful. The `require` function checks for a specific condition to be true. If this is not the case, the function will be reverted and all the gas goes back to the sender, which can be evaluated in Solidity using `msg.sender`. In

order to prevent over- or underflow errors when performing additions and subtractions, a small library **SafeMath** (program 4.1) is being utilized. If everything worked well, the **Transfer** event will be emitted.

```

1 function transfer(address _to, uint256 _value) public returns (bool success) {
2   require(_value <= balances[msg.sender]);
3   balances[msg.sender] = balances[msg.sender].sub(_value);
4   balances[_to] = balances[_to].add(_value);
5   emit Transfer(msg.sender, _to, _value);
6   return true;
7 }

```

The function **transferFrom** is very similar to **transfer** with the only difference that the caller is not the one who sends the tokens. This is usually not a physical person but rather a smart contract of a decentralized application. Tokens can be sent only on behalf of someone else if the sender accepts the amount first. All funds that are already released for transactions are stored in another mapping specifying for each address which addresses are enabled to send what amount.

```

1 mapping(address => mapping(address => uint256)) private allowed;
2
3 function transferFrom(address _from, address _to, uint256 _value) public returns (
4   bool success) {
5   require(_value <= balances[_from]);
6   require(_value <= allowed[_from][msg.sender]);
7   balances[_from] = balances[_from].sub(_value);
8   allowed[_from][msg.sender] = allowed[_from][msg.sender].sub(_value);
9   balances[_to] = balances[_to].add(_value);
10  emit Transfer(_from, _to, _value);
11  return true;
12 }

```

If the caller is not the sender of a transaction, funds need to be approved first. That is exactly what the function **approve** does. It authorizes the spender to spend a specific amount of the tokens of the address which called the function. Old funds that are already authorized will not be accumulated but overwritten. This makes it possible to revert an approval by calling the **approve** function another time with the value of zero. This function emits the **Approval** event.

```

1 function approve(address _spender, uint256 _value) public returns (bool success) {
2   allowed[msg.sender][_spender] = _value;
3   emit Approval(msg.sender, _spender, _value);
4   return true;
5 }

```

Because the data structure holding the allowed funds is usually private, a public interface is needed to find out how much remaining tokens can be spent on behalf of someone else. The function **allowance** takes two parameters and returns the amount of tokens that the spender is authorized to spend from the owner's address.

```

1 function allowance(address _owner, address _spender) public view returns (uint256
2   remaining) {
3   return allowed[_owner][_spender];
4 }

```

The events **Transfer** and **Approval** play an important role in the interface of an ERC20 token. The **Transfer** event is emitted every time when a transaction is about

to be executed. The `Approval` event is being called every time some tokens need to be approved. Both events trigger some sort of action in the linked wallet application of the user. For example, the browser wallet MetaMask opens a popup window asking the user for confirmation, once an event is emitted.

```
1 event Transfer(address indexed _from, address indexed _to, uint256 _value);
2 event Approval(address indexed _owner, address indexed _spender, uint256 _value);
```

4.4 Liquidity Pool Implementation

There are a few different approaches on how to implement a lending functionality. One possibility would be the utilization of collateral like it is done by Maker [36]. However, this section will describe how to implement a lending application using liquidity pools as it builds the foundation for decentralized exchanges such as Uniswap [40].

As described in section 3.5.2, liquidity pools always need to have a pair of digital assets. As we are building on top of the Ethereum network, we are limited to Ether and Ethereum tokens. In this example, the liquidity pool holds Ether and an ERC20 token called *CherryToken*, as implemented in section 4.3. Besides the token reference `cherryToken`, the smart contract needs to keep track of the owner, rewards that are ready to be distributed to the liquidity providers and a mapping where the balances of the liquidity providers are stored. In order to create a single measurement for the pool pair, a virtual liquidity token is being created. It can be calculated by multiplying both assets and computing the square root. Because dealing with floating point numbers is difficult in Solidity and this is a redundant information which can be calculated, it is done by the client and not directly on the blockchain. This is a common practice to store only absolute necessary data directly on the blockchain in order to save memory space and gas fees.

```
1 address private _owner;
2 CherryToken private _cherryToken;
3
4 uint256 private _availableRewards = 0;
5
6 mapping(address => uint256) private _liquidityTokenBalances;
```

On Ethereum, not only externally owned accounts (EOAs) can receive Ether. Smart contracts are able to receive payments by adding the `payable` modifier to a function. Most contracts have an empty payable function, called fallback function, to prevent the loss of Ether which is being sent to the contract without calling a specific function. The getter functions `getEthBalance`, `getCtnBalance`, `getAvailableRewards` and `getLiquidityTokenBalance` provide the frontend with necessary information to implement a good user experience and, again, to keep as much calculation off-chain.

```
1 function () external payable {}
2
3 function getEthBalance() public view returns (uint256 balance) {
4     return address(this).balance;
5 }
6
7 function getCtnBalance() public view returns (uint256 balance) {
8     return _cherryToken.balanceOf(address(this));
```

```

 9 }
10
11 function getAvailableRewards() public view returns (uint256 rewards) {
12     return _availableRewards;
13 }
14
15 function getLiquidityTokenBalance(address owner) public view returns (uint256 funds)
16     {
17     return _liquidityTokenBalances[owner];
18 }

```

Adding liquidity works a little bit different for Ether and an ERC20 token, as Ether is a fundamental feature of Solidity and an ERC20 token is managed by the smart contract itself. Like the fallback function, which just adds anonymized Ether to the liquidity pool, `addLiquidity` detects the address of the sender in order to issue liquidity tokens. After a few integrity checks, the token liquidity is transferred to the contract and the balance of the liquidity provider gets updated.

```

1 function addLiquidity(uint256 ctnAmount, uint256 liquidityTokenAmount) public
  payable {
2     require(msg.value > 0, "ETH amount cannot be 0");
3     require(_cherryToken.balanceOf(msg.sender) >= ctnAmount, "not enough CTN funds");
4     _cherryToken.transferFrom(msg.sender, address(this), ctnAmount);
5     _liquidityTokenBalances[msg.sender] = _liquidityTokenBalances[msg.sender].add(
6         liquidityTokenAmount);
7 }

```

Removing liquidity from the pool means withdrawing the loaned funds and receiving an additional reward proportional to the pool share. As the reward calculation is quite computational intensive and Solidity is not good at managing divisions and percentages, the reward will be computed off-chain as well. Because this is not managed by the function `removeLiquidity` itself, there are quite a few `require` statements necessary, to detect potential errors. The rewards are created using the `mint` function from the token and paid out proportional to the share of liquidity holdings of the investor.

```

1 function removeLiquidity(uint256 ethAmount, uint256 ctnAmount, uint256
  liquidityTokenAmount, uint256 totalLiquidityTokens) public {
2     require(ethAmount > 0, "eth amount cannot be 0");
3     require(ctnAmount > 0, "ctn amount cannot be 0");
4     require(getEthBalance() >= ethAmount, "not enough ETH in pool");
5     require(getCtnBalance() >= ctnAmount, "not enough CTN in pool");
6     require(_liquidityTokenBalances[msg.sender] > liquidityTokenAmount, "not enough
  liquidity tokens");
7     uint256 ctnRewards = _availableRewards.mul(liquidityTokenAmount.div(
  totalLiquidityTokens));
8     _availableRewards = _availableRewards.sub(ctnRewards);
9     _cherryToken.mint(msg.sender, ctnRewards);
10    _cherryToken.transfer(msg.sender, ctnAmount);
11    msg.sender.transfer(ethAmount);
12    _liquidityTokenBalances[msg.sender] = _liquidityTokenBalances[msg.sender].sub(
  liquidityTokenAmount);
13 }

```

4.5 Exchange Implementation

A truly decentralized exchange is dependent of its own liquidity pools. While the concept of a liquidity pool is quite complex, the actual implementation of the exchange is a lot easier. It can be implemented in just two functions. One function for each exchange direction. Both of these functions follow the same structure: Receiving funds, calculating fees, deducting fees and sending funds.

When exchanging from Ether to an ERC20 token, the function needs to be payable again. The incoming payment is accessible in the function by using the `value` property of the message object. After checking if the payment is not empty, the function calculates the fees which will be distributed to the liquidity providers. This calculation can be either fixed or dynamically. On more sophisticated exchanges, the percentage of the fees uses an approximation of an exponential curve. If the liquidity pool is very small, the fees are really high¹. As the pool gains more depth, the fees decrease into small fractions of percentages but never will be zero.

All numeric asset representations are stored in wei, the smallest unit on the Ethereum platform. One Ether is equivalent to one quintillion wei, which is a number with 19 digits.

```
1 function swapEthToCtn(uint256 ctnOutput, uint256 fees) public payable {
2   require(msg.value > 0, "amount cannot be 0");
3   _addFees(fees);
4   require(getCtnBalance() >= ctnOutput, "not enough funds available");
5   _cherryToken.transfer(msg.sender, ctnOutput);
6 }
```

After the fees are calculated, the function checks, if the smart contract has enough funds to pay the user back. If this is not the case, the transaction aborts and the money will be transferred back to the user. If there are still enough funds available, the amount deducted by the fees will be paid back using the `transfer` function from the ERC token standard.

The reversed function is not equipped with the `payable` modifier, as the incoming funds are an ERC20 token. Because the funds of the ERC20 tokens are managed by the token itself, the exchange contract has no direct access to it. If a user sends token funds to the address of the exchange, the exchange has no way to get notified about it. Therefore, the token transaction will take place in the contract function on behalf of the user. All the user has to do before calling the exchange function, is to grant the exchange permission to spend the funds on his behalf using the `approve` function from the ERC20 token standard.

Instead of just paying the desired amount to the exchange, the amount needs to be passed to the function as a parameter. As before, the amount is being checked if it is empty to save the gas which would be wasted when executing a swap with zero funds. After that, the exchange rate and fees will be calculated. After another safety check, the tokens will be finally sent from the user to the exchange and the exchange sends the Ether amount back to the user using the built-in `transfer` function.

```
1 function swapCtnToEth(uint256 ctnInput, uint ethOutput, uint256 fees) public {
2   require(ctnInput > 0, "amount cannot be 0");
3   _addFees(fees);
```

¹up to 50% and more

```

4   require(getEthBalance() >= ethOutput, "not enough funds available");
5   _cherryToken.transferFrom(msg.sender, address(this), ctnInput);
6   msg.sender.transfer(ethOutput);
7 }

```

4.6 Testing Smart Contracts

Testing plays a crucial part in the development of decentralized applications. Once deployed, the smart contracts are immutable on the blockchain. Fixing bugs after the initial deployment can be expensive and tedious, as the affected smart contracts need to be re-deployed and might be incompatible with the previous version. When working with financial assets of other people, it is even more vital, to minimize the number of bugs. Therefore, most applications use one or more testing frameworks and use a three-phase deployment cycle, as described in more detail in section 4.8.

As a developer, you have the choice to write tests natively in Solidity, or using an arbitrary testing framework based on JavaScript. When working with Truffle, the testing framework Mocha [22] and the assertion library Chai [17] are preinstalled as the default. All tests that are in the testing subdirectory, will be executed using the command `truffle test`.

At the top of each test file, the artifacts of the used smart contracts, as well as Chai, need to be imported:

```

1 const CherryPool = artifacts.require('CherryPool');
2 const CherryToken = artifacts.require('CherryToken');
3
4 require('chai')
5   .use(require('chai-as-promised'))
6   .should();

```

The specific tests are to be implemented in a contract block, which takes two parameters: The name of the tested contract and an array of test accounts as parameters. As the first account is used to deploy the contract, it is labelled as the contract owner.

```

1 contract('CherryPool', ([owner, user]) => {
2   // tests are here
3 }

```

While the contract compilation takes place before the tests are being executed, the smart contracts still need to be instantiated. This can be done using a `before` block, where all initialization tasks for the tests take place:

```

1 before(async () => {
2   // load contracts
3   cherryToken = await CherryToken.new();
4   cherryPool = await CherryPool.new(cherryToken.address);
5   await cherryToken.transfer(cherryPool.address, '1000000000000000000000');
6   await web3.eth.sendTransaction({ from: owner, to: cherryPool.address, value:
       web3.utils.toWei('1') });
7 });

```

The actual tests are grouped together to tests that are similar to each other or that test the same functionality into `describe` and `it` blocks for the concrete unit

tests. The estimated results of the test can be checked with assertions or by appending `should.be.rejected` to a function call.

```

1 describe('CherryPool deployment', async () => {
2   it('has no rewards yet', async () => {
3     const rewards = await cherryPool.getAvailableRewards();
4     assert.equal(web3.utils.fromWei(rewards), 0);
5   });
6
7   it('has initial Ether pool', async () => {
8     const ethBalance = await cherryPool.getEthBalance();
9     assert.equal(web3.utils.fromWei(ethBalance), 1);
10  });
11
12  it('has initial CherryToken pool', async () => {
13    const ctnBalance = await cherryPool.getCtnBalance();
14    assert.equal(web3.utils.fromWei(ctnBalance), 1000);
15  });
16 });

```

4.7 Connecting Smart Contracts to Web Applications

The connection between the web frontend and the smart contracts is done using the JavaScript framework Web3 [28], which is maintained by the Ethereum core team. It contains a lot of useful utility functions as seen before such as `web3.eth.sendTransaction()` for making transaction or `web3.utils.fromWei()` to convert a number from wei into Ether.

When the user wants to connect to a wallet, Web3 needs to determine the network provider first. Once detected, a new instance of Web3 will be instantiated and is ready for use.

```

1 function loadWeb3() {
2   if ((window as any).ethereum) {
3     (window as any).web3 = new Web3((window as any).ethereum);
4     (window as any).ethereum.enable().then();
5     return true;
6   } else if ((window as any).web3) {
7     (window as any).web3 = new Web3((window as any).web3.currentProvider);
8     return true;
9   }
10
11   window.alert('Non-Ethereum browser detected. You should consider trying MetaMask!');
12   return false;
13 }

```

The accounts array can be accessed via `web3.eth.getAccounts()` and the network id is retrieved using `web3.net.getId()`. Each contract can be instantiated afterwards based on the compilation artifacts. In addition to the tests, the correct network needs to be selected first.

```

1 const cherryPoolData = (CherryPool as any).networks[networkId];
2 if (cherryPoolData) {
3   cherryPool = new web3.eth.Contract((CherryPool as any).abi, cherryPoolData.address);
4 }

```



```
4 } else {  
5   alert('CherryPool contract not deployed to detected network!');  
6   return null;  
7 }
```

After these steps, the API is ready to be used by any UI component. Smart contract functions can be invoked using `call()` or `send()`. The function `call()` is used if the function does not alter the state of the blockchain. A typical use case for that is retrieving an account balance:

```
1 account.cherryToken.methods.balanceOf(account.address).call().then((ctnBalance:  
    number) => {  
2   setCherryTokenBalance(web3.utils.fromWei(ctnBalance.toString()));  
3 });
```

If a feature needs to alter the state of the blockchain, the function `send()` needs to be used. As every state modification costs gas, this will usually lead to a confirmation dialog provided by the wallet of the user. Typical use cases for that are adding liquidity, exchanging assets or approving funds for other accounts:

```
1 account.cherryToken.methods.approve(  
2   account.cherryPool._address,  
3   web3.utils.toWei(ctnToSupply.toString())  
4 ).send({ from: account.address }).then();
```

4.8 Smart Contract Deployment

A typical release cycle of a blockchain application consists of three testing environments. At first, the smart contracts will be deployed to a local blockchain such as Ganache [20]. That is the environment where most of the testing will take place. Once the application is quite stable, it will be migrated to a public test network. These networks provide more realistic conditions, as it behaves more like the main network. Assets on the test networks have no value, which is really difficult to implement in reality. In order to prevent a value increase, the test networks need to be redeployed from time to time.

The most important test networks on Ethereum are *Ropsten*, *Kovan* (using the Aura consensus protocol) and *Rinkeby* (using the Clique consensus protocol). These consensus protocols are specifically designed to prevent the creation of value, which is exactly the opposite what the consensus protocols of proof-of-work and proof-of-stake are trying to achieve.

Chapter 5

Closing Remarks

5.1 Criticism

The previous chapters addressed the advantages, benefits and possibilities of decentralized finance. But as always, there is no free lunch. Even decentralized finance is not immune against criticism. This section provides a brief overview of a few of the most relevant issues coming from detractors as well of a data-based evaluation whether they are valid or not.

Cryptocurrencies like Bitcoin that are using a proof-of-work algorithm in order to achieve consensus have an extremely high power consumption. Figure 5.1 shows that the Bitcoin network alone has a yearly electrical energy use of 77.78 TWh, just passing Austria with 74.1 TWh in October 2020. While this are numbers, that cannot be blanded, it is important to keep in mind that the mining algorithm on Bitcoin is jointly responsible for its high value and is vital for the network, in order to operate in a proper and secure way.

A solution would be a proof-of-stake algorithm which is the choice of the majority of newer coins. Ethereum is already moving to proof-of-stake as part of its update Ethereum 2.0 but that is not happening on Bitcoin anytime soon. At least 73% of Bitcoin's electricity comes from renewable energy sources, as investigated by researches from *CoinShares* [5, p. 8]. This has economical reasons. In an open market economy, Bitcoin miners will always move to those places where energy is cheaper than anywhere else. The majority of Bitcoin mining is done in Sichuan, China because the hydroelectricity is very cheap in the rainy season. But there are also other attractive locations such as Iceland and its eco-friendly geothermal energy sources.

Another popular objection is the high use of cryptocurrencies for criminality and money laundering. While there is obviously no data about the actual numbers of value used illegally, the majority of illicit assets still take place in fiat currencies. It is only logical, that cryptocurrencies are the preferred choice to be used on dark net markets. But that is not a problem of the cryptocurrencies themselves, as like any technology, they are just tools which can help to achieve either good or bad things. Shifting the focus back to decentralized finance in particular, the number of illegal activities is plummeting, as there are even more privacy focused coins like *Zcash*, *Monero* and *Dash* [7, p. 13].

As of the arise of a lot new projects in the area of decentralized finance, some of




Figure 5.1: Bitcoin’s Energy Consumption Compared to Countries [30].

them are exposed to ponzi schemes and similar frauds, as fairly new projects usually are very centralized and only dispense control after time. There is no real solution for that, except encouraging people to do their own research beforehand and by following the principle of only investing in what you understand. But even if the founders have no bad intentions, it is still possible that users loose money due to errors in the code. This is specifically a real problem on Ethereum, as of its turing-completeness increases the risk of unconsciously adding bugs to the code. *DeFiChain* [31], a Bitcoin fork specifically adapted for decentralized financial use cases, tries to solve that problem by restricting the interaction with the blockchain by only using few commands.

5.2 Risks

Investing in decentralized financial products has certain risks. As already mentioned in section 5.1, projects on the Ethereum network are exposed to a high smart contract risk, due to the open design of its programming languages. Coding errors may result in loss of funds, which are lost forever. However, the smart contract risk on Bitcoin is negligibly low as its commands are being executed on a simple stack machine.

The cryptocurrency market as a whole is very volatile. While this is neither a good nor a bad thing, it definitely increases the overall risk. Especially for new investors are exposed to make regrettable emotional decisions without doing fact-based research. Generally speaking, high volatility is good for people with a long time horizon and who are focusing on capital gains. If people have a short time horizon or aim for regular cashflow, they should rather avoid markets with high volatility.

A very insidious type of risk are those, which cannot be predicted. This could be

scams, hacks or other so called black swan events. A *black swan* is by definition an event which is extremely unlikely to happen, not predictable and it would cause potentially severe consequences. Examples for black swan events on Bitcoin would be an evil nature of its founder Satoshi Nakamoto or a severe error in the fundamental cryptographic algorithms.

Decentralization comes with many advantages in certain areas. But as there is a lot of freedom, people can do either good or bad things on such an ecosystem. When thinking about tokenization, it will be possible soon to bring real world assets on a blockchain. For example, Starbucks is already about to move their supply chain on a blockchain, where it will be completely transparent for the customer to see where the coffee beans came from and which middlemen were involved. By scanning the QR code on the cup it will be even possible to tip the farmer on the coffee plantation, instead of just the cashier in the store. But what if people will be tokenized too? Well, that is already reality. In 2020, basketball player Spencer Dinwiddie tokenized his NBA contract on the Ethereum blockchain, which raised \$1.3 million [37]. While this is still quite harmless right now, there are also some dystopian scenarios like betting on a person's death.

5.3 Prospective Impact

It is inevitable that decentralized finance will continue rising in popularity in the next few years. Every new affair of a central authority proves that again and again. We have already seen that many times in late 2020 and early 2021. Be it the tremendously high amount of new fiat currencies being printed, blocking certain content and users on social media or by prohibiting retail investors to trade certain securities in the *GameStop short squeeze*. It is all just another nail in the coffin of the traditional banking system and other centralized organizations. But the technology is not ready yet for the mass adoption and using digital decentralized currencies on a day to day basis will certainly not be coming in the next few years.

References

Literature

- [1] Andreas Antonopoulos. *Mastering Bitcoin: Programming the Open Blockchain*. 2nd ed. Sebastopol: O'Reilly Media, Inc., 2017 (cit. on p. 9).
- [2] Andreas Antonopoulos and Gavin Wood. *Mastering Ethereum: Building Smart Contracts and Dapps*. Sebastopol: O'Reilly Media, Inc., 2018 (cit. on pp. 10, 16, 18).
- [3] Adam Back. *Hashcash - A Denial of Service Counter-Measure*. Aug. 2002. URL: <http://www.hashcash.org/papers/hashcash.pdf> (cit. on p. 9).
- [4] Badr Bellaj, Richard Horrocks, and Xun Wu. *Blockchain By Example*. Birmingham: Packt Publishing Ltd., 2018 (cit. on p. 18).
- [5] Christopher Bendiksen and Samuel Gibbons. *The Bitcoin Mining Network - Trends, Composition, Average Creation Cost, Electricity Consumption & Sources*. Tech. rep. London, UK: CoinShares Research, Dec. 2019. URL: <https://coinshares.com/assets/resources/Research/bitcoin-mining-network-december-2019.pdf> (cit. on p. 26).
- [6] Daniel R. L. Brown. *SEC 2: Recommended Elliptic Curve Domain Parameters*. Standards for Efficient Cryptography. Jan. 27, 2010. URL: <https://www.secg.org/sec2-v2.pdf> (cit. on p. 4).
- [7] Yaya J. Fanusie and Tmo Robinson. *Bitcoin Laundering: An Analysis of Illicit Flows into Digital Currency Services*. Tech. rep. Washington D.C., USA: Foundation for Defense of Democracies, Jan. 2018. URL: https://s3.us-east-2.amazonaws.com/defenddemocracy/uploads/documents/MEMO_Bitcoin_Laundering.pdf (cit. on p. 26).
- [8] Brenn Hill, Chopra Samanyu, and Valencourt Paul. *Blockchain Quick Reference*. Birmingham: Packt Publishing Ltd., 2018 (cit. on p. 10).
- [9] Satoshi Nakamoto. *Bitcoin: A Peer-to-Peer Electronic Cash System*. Oct. 2008. URL: <https://bitcoin.org/bitcoin.pdf> (cit. on p. 9).
- [10] Klaus Schmeh. *Kryptografie: Verfahren, Protokolle, Infrastrukturen*. Heidelberg: dpunkt.verlag GmbH, 2007 (cit. on pp. 3, 4).
- [11] *Solidity Documentation*. Version 0.7.5. Oct. 2020. URL: <https://docs.soliditylang.org/> (cit. on p. 16).

- [12] Melanie Swan. *Blockchain: Blueprint for a New Economy*. Sebastopol: O'Reilly and Associates, 2015 (cit. on p. 9).
- [13] Fabian Vogelsteller and Buterin Vitalik. *ERC-20 Token Standard*. A standard interface for tokens. Mar. 8, 2019. URL: <https://github.com/ethereum/EIPs/blob/master/EIPS/eip-20.md> (cit. on p. 18).
- [14] Gavin Wood. *Ethereum: A Secure Decentralised Generalised Transaction Ledger*. Version 3e2c089. Sept. 2020. URL: <https://ethereum.github.io/yellowpaper/paper.pdf> (cit. on p. 10).
- [15] Xun Wu and Weimin Sun. *Blockchain Quick Start Guide*. Birmingham: Packt Publishing Ltd., 2018 (cit. on p. 13).

Software

- [16] *Ampleforth*. URL: <https://github.com/ampleforth/uFragments> (visited on 01/10/2021) (cit. on p. 11).
- [17] *Chai*. URL: <https://github.com/chaijs/chai> (visited on 11/26/2020) (cit. on pp. 17, 23).
- [18] *Chainlink: Blockchain Oracles for connected Smart Contracts*. URL: <https://github.com/smartcontractkit/chainlink> (visited on 01/10/2021) (cit. on p. 8).
- [19] *CherryPool*. URL: <https://github.com/johannesdominik/cherrypool> (visited on 11/26/2020) (cit. on p. 16).
- [20] *Ganache*. URL: <https://github.com/trufflesuite/ganache> (visited on 11/26/2020) (cit. on pp. 17, 25).
- [21] *Go Ethereum*. URL: <https://github.com/ethereum/go-ethereum> (visited on 01/31/2021) (cit. on p. 5).
- [22] *Mocha*. URL: <https://github.com/mochajs/mocha> (visited on 11/26/2020) (cit. on pp. 17, 23).
- [23] *Parity*. URL: <https://github.com/openethereum/parity-ethereum> (visited on 01/30/2021) (cit. on p. 5).
- [24] *React*. URL: <https://github.com/facebook/react/> (visited on 11/26/2020) (cit. on p. 17).
- [25] *The Solidity Contract-Oriented Programming Language*. URL: <https://github.com/ethereum/solidity> (visited on 11/26/2020) (cit. on p. 16).
- [26] *Truffle*. URL: <https://github.com/trufflesuite/truffle> (visited on 11/26/2020) (cit. on p. 16).
- [27] *TypeScript*. URL: <https://github.com/Microsoft/TypeScript> (visited on 01/09/2021) (cit. on p. 17).
- [28] *web3.js - Ethereum JavaScript API*. URL: <https://github.com/ethereum/web3.js> (visited on 11/26/2020) (cit. on pp. 17, 24).

Online sources

- [29] *Bitcoin Cash*. Nov. 2020. URL: <https://www.bitcoincash.org/> (visited on 11/24/2020) (cit. on p. 10).
- [30] *Bitcoin Energy Consumption Index*. Jan. 2021. URL: <https://digiconomist.net/bitcoin-energy-consumption> (visited on 01/29/2021) (cit. on p. 27).
- [31] *Decentralized finance enabled on Bitcoin*. Jan. 2021. URL: <https://defichain.com/> (visited on 01/29/2021) (cit. on pp. 7, 27).
- [32] *DeFi Pulse*. Jan. 2021. URL: <https://defipulse.com/> (visited on 01/29/2021) (cit. on p. 7).
- [33] *Ethereum Casper*. Sept. 2018. URL: <https://twitter.com/i/events/1036281460704112645> (visited on 11/24/2020) (cit. on p. 10).
- [34] *Gemini Dollar*. Jan. 2021. URL: <https://www.gemini.com/dollar> (visited on 01/10/2021) (cit. on p. 11).
- [35] *Liquidity Mining, risks and exchange function on a DEX explained*. Nov. 2020. URL: <https://julianhosp.com/liquidity-mining-risks-and-exchange-function-on-a-dex-explained/> (visited on 11/24/2020) (cit. on p. 14).
- [36] *MakerDAO: An Unbiased Global Financial System*. Jan. 2021. URL: <https://makerdao.com/en/> (visited on 01/10/2021) (cit. on pp. 11, 12, 20).
- [37] Jamie Redman. *NBA Point Guard Spencer Dinwiddie's Tokenized Contract Raises \$1.3 Million*. July 2020. URL: <https://news.bitcoin.com/nba-point-guard-spencer-dinwiddies-tokenized-contract-raises-1-3-million/> (visited on 01/30/2021) (cit. on p. 28).
- [38] *Tether: Stable digital cash on the Blockchain*. Jan. 2021. URL: <https://tether.to/> (visited on 01/10/2021) (cit. on p. 11).
- [39] *Top DeFi Tokens by Market Capitalization*. Jan. 2021. URL: <https://coinmarketcap.com/defi/> (visited on 01/09/2021) (cit. on p. 8).
- [40] *Uniswap Pools*. Nov. 2020. URL: <https://uniswap.org/docs/v2/core-concepts/pools/> (visited on 11/24/2020) (cit. on pp. 13, 20).
- [41] *USD Coin (USDC) Stablecoin*. Jan. 2021. URL: <https://www.circle.com/en/usdc> (visited on 01/10/2021) (cit. on p. 11).

Check Final Print Size

— Check final print size! —



— Remove this page after printing! —